

**Hewlett Packard
Enterprise**

Chapel's Performance Tracking System

Jade Abraham
July 10th, 2025

What is Chapel?

Chapel: A modern parallel programming language

- portable & scalable
- open-source & collaborative



Goals:

- Support general parallel programming
- Make parallel programming at scale far more productive

chapel-lang.org



Overview

- Goals of performance tracking
- Configurations we test
 - Networks
 - Single node vs Multi node
 - Backend Compilers
 - Key User Applications
- Key Features
 - Graphs
 - Data storage
 - Custom benchmarking
 - Playgrounds
- Areas for improvement



Goals

- Track Benchmarks Written in Chapel
 - Regressions and improvements
 - Need ability to triage and learn from regressions
 - Also need to understand what causes improvements, can that drive future improvements?
- Track Key User Applications
 - We support users, we should find performance regressions before they do
 - Arkouda, CHAMPS, ChOp, miniBUDE
- Track Compiler Performance
 - Prevent compilation time/memory consumption from regressing
- Transparency
 - All our perf results are publicly available, the good and the bad
 - <https://chapel-lang.org/perf-nightly.html>



Current Configurations (Non-exhaustive list)

- Single Node (24 cores and 256 GB RAM)
 - Default Configuration
 - C Backend – GNU and clang
 - Compiler Performance – Default and Optimized
 - Arkouda (User Application)
- HPE Slingshot runs (200 Gb/s Slingshot 11, 64 cores and 512 GB RAM)
 - 16 Node
 - 1 Node
- InfiniBand runs (200 Gb/s HDR, 72 cores and 512 GB RAM)
 - 16 Node
 - 16 Node - Arkouda (User Application)
 - 16 Node - CHAMPS (User Application)


Current Configurations (Non-exhaustive list)

- Cray XC (node's have 48 cores and 192 GB RAM)
 - 16 Node
 - 1 Node
 - 16 Node - Arkouda (User Application)
- Single Node GPU
 - Nvidia A100 (4 GPUs, 64 CPUs)
 - AMD MI250x (4 GPUs, 64 CPUs)



Graphs

- Graphs are described in .graph files
- A single graph file can pull in multiple data files
- Generation process
 - Python script to parse all graph files + data + annotations into graphdata.js
 - Custom web scripting to generate graphs (dygraph)
- Annotations
 - YAML file of annotations for graphs
 - Attach annotations to entire configurations or individual tests



<https://chapel-lang.org/perf/16-node-apollo-hdr?graphs=baleindexgatherperfmbsslocale,baleaggregatedindexgatherperfmbsslocale,baleindexgathercompileroptvariantsperfmbsspernode,baleindexgathertextimplicitoptvariantsperfmbsspernode,baleindexgathertimesec>

Live Graphs!

- Single Node: <https://chapel-lang.org/perf/chapcs/>
 - Shootout
 - Perf Tracking – Last Two Weeks
- 16 Node InfiniBand: <https://chapel-lang.org/perf/16-node-apollo-hdr/>
 - Selected graphs

Infrastructure

- Chapel's test system supports correctness and performance tests
- Each nightly job is configured by
 - Public scripts in the Chapel repo
 - Internal Jenkins configuration
- Jenkins jobs run nightly on hardware hosted by HPE
- Data
 - Data is saved into tab separated .dat files
 - Sync'd to a GitHub repository for historical records
- Graphs
 - Graphs are regenerated each night by the Jenkins job
 - Separate Jenkins job collects graphs and automatically uploads them to webhost

```
util > cron > $ test-perf.chapcs.bash
1  #!/usr/bin/env bash
2  #
3  # Run performance tests on a chapcs machine with default configuration.
4
5 UTIL_CRON_DIR=$(cd $(dirname ${BASH_SOURCE[0]}) ; pwd)
6
7 export CHPL_TEST_PERF_CONFIG_NAME='chapcs'
8
9 source $UTIL_CRON_DIR/common-perf.bash
10
11 export CHPL_LAUNCHER=none
12
13 export CHPL_NIGHTLY_TEST_CONFIG_NAME="perf.chapcs"
14
15 $UTIL_CRON_DIR/nightly -cron -performance -releasePerformance -numtrials 5 -startdate
09/10/15
16 |
```

Writing Performance Tests

- Most performance tests do not require any custom scripting
- Each test is a Chapel file with some extra config files
 - `perfcompopts`
 - Compilation options passed to the Chapel compiler
 - `perfexecopts`
 - Options passed to the compiled executable
 - `perfkeys`
 - Defines characteristics to extract from test for data files
 - Support for verifying correct output from test in same perf test
- Built in support for testing other languages for comparison, mostly C
 - Can test just about anything via custom scripting
- We can also plug in arbitrary scripts to test larger/more complex projects
 - This is the strategy for Arkouda and CHAMPS



Example Test: test/npb/ep/mcahir/ep.chpl

- ep.perfexecopts: defines multiple variations on the same benchmark

```
test > npb > ep > mcahir >  ≡ ep.perfexecopts
1  --probClass=S # ep.S
2  --probClass=W # ep.W
3  --probClass=A # ep.A
4  --probClass=B # ep.B
5  #--probClass=C # ep.C
6  #--probClass=D # ep.D
```

- ep.perfkeys: for each variation run, defines the performance keys to look for

```
test > npb > ep > mcahir >  ≡ ep.perfkeys
1  Time in seconds =
2  Mop/s total      =
3  Mop/s/process    =
4  Verification     =
5  verify:-1: ^ Verification .*SUCCESSFUL$
```

Example Test: test/npb/ep/mcahir/ep.chpl

- ep.S.dat: one of the data files generated

≡ ep.S.dat					
1	# Date	Time in seconds =	Mop/s total	=	Mop/s/process
2	09/10/15	0.08683	386.438	=	386.438 SUCCESSFUL
3	09/10/15	0.092351	363.336	=	363.336 SUCCESSFUL
4	09/10/15	0.090567	370.493	=	370.493 SUCCESSFUL
5	09/10/15	0.095063	352.97	=	352.97 SUCCESSFUL
6	09/10/15	0.094059	356.738	=	356.738 SUCCESSFUL
7	09/11/15	0.088565	378.868	=	378.868 SUCCESSFUL
8	09/11/15	0.09379	357.761	=	357.761 SUCCESSFUL
9	09/11/15	0.089163	376.327	=	376.327 SUCCESSFUL
10	09/11/15	0.088058	381.049	=	381.049 SUCCESSFUL
11	09/11/15	0.088211	380.388	=	380.388 SUCCESSFUL
12	09/12/15	0.083856	400.143	=	400.143 SUCCESSFUL



Example Test: test/npb/ep/mcahir/ep.chpl

- ep.graph: defines how the data is presented

```
test > npb > ep >  ≡ ep.graph
 1  perfkeys: Time in seconds =, Time in seconds =
 2  files: ep.S.dat, ep.S.maxLogical.dat
 3  graphkeys: runtime - S, MAX_LOGICAL - S
 4  graphtitle: NAS Parallel Benchmarks: EP timings - size S
 5  ylabel: Time (seconds)
 6
 7  perfkeys: Time in seconds =, Time in seconds =
 8  files: ep.W.dat, ep.W.maxLogical.dat
 9  graphkeys: runtime - W, MAX_LOGICAL - W
10  graphtitle: NAS Parallel Benchmarks: EP timings - size W
11  ylabel: Time (seconds)
12
13  perfkeys: Time in seconds =, Time in seconds =
14  files: ep.A.dat, ep.A.maxLogical.dat
15  graphkeys: runtime - A, MAX_LOGICAL - A
16  graphtitle: NAS Parallel Benchmarks: EP timings - size A
17  ylabel: Time (seconds)
18
```


Example Test: test/npb/ep/mcahir/ep.chpl

Chapel Performance Graphs for chapcs



Playgrounds

- We have two playgrounds
 - Single node
 - 16 node InfiniBand
- Allows developers to test new features or dependencies without disrupting main
- Example: <https://chapel-lang.org/perf/chapcs/llvm-16/?graphs=hpccratime>



```
11 util/cron/test-perf.chapcs.playground.bash □
20   20
21   21     export CHPL_NIGHTLY_TEST_CONFIG_NAME="perf.chapcs.playground"
22   22
23   23     #
24   24     # 0) Update the comment that follows this block to describe what's being tested
25   25     # 1) Update GITHUB_USER to the GitHub username owning the branch
26   26     # 2) Update GITHUB_BRANCH to the name of the GitHub branch under that user
27   27     # 3) Update SHORT_NAME to be a short name that will be used in graph generation
28   28     # 4) Update START_DATE to be today, using the format mm/dd/yy
29   29     #
30   30
31 - GITHUB_USER=chapel-lang
32 - GITHUB_BRANCH=main
33 - SHORT_NAME=libmvec
34 - START_DATE=04/16/25
31 + GITHUB_USER=jabraham17
32 + GITHUB_BRANCH=llvm-20
33 + SHORT_NAME=llvm20
34 + START_DATE=04/30/25
```


Triage

- Weekly
 - Several members of the Chapel team meet to review the last week's performance data
 - Look at key configs and user applications for regressions/improvements
 - Assign people to reproduce and diagnose changes
- Before every release
 - The Chapel team goes through the performance data since the previous release (~3 months)
 - Look at all configs and check for anything that was missed




Recent Case Study #1

- Arkouda developer added new support for multi-dimensional arrays
- Array creation regressed
 - We could immediately tell the problem was user created by testing both Chapel nightly and release
 - <https://chapel-lang.org/perf/arkouda/chapcs/?startdate=2023/11/04&enddate=2023/12/21&graphs=streamperformance>
- Developer was able to diagnose the root cause and fix the regression



Recent Case Study #2

- We upgraded our nightly system to use LLVM 20
- Bitops comparison test regressed
 - The test (written in C) is a microbenchmark of bitwise operations in the Chapel runtime
 - <https://chapel-lang.org/perf/chapcs/?startdate=2025/05/02&enddate=2025/07/03&graphs=bitopspopcountc,bitopsclzc,bitopscszc>
- Found a LLVM bug and a bug in the performance test
 - <https://github.com/llvm/llvm-project/issues/142042>
 - <https://github.com/chapel-lang/chapel/pull/27308>



Areas of Improvements

- New Features
 - Add git SHA next to the date in the graphs
 - Add information on what file and command line options create a given line on a graph
 - Split data files on more than just a per-file basis
 - Adding/removing a perf key to an existing file today requires editing the historical data
 - Ability to setup perf tests from a single configuration file, rather than multiple files
 - Arrows to show where a regression is
 - Some graphs are “up is better”, some are “down is better”. This routinely trips people up.
- Modernizing our testing infrastructure
 - We have infrastructure that works for us and provides results, but it is built on aging software
 - Many scripts are 20+ years old and few people know how they work
 - We rely heavily on bespoke solutions rather than standard tools
 - Managing Jenkins and system configurations is currently a big time sink

Thank you

<https://chapel-lang.org>
@ChapelLanguage

