

Comparing Distributed-Memory Programming Frameworks with Radix Sort

Michael Ferguson (HPE), Ryan Friese (PNNL), Shreyas Khandekar (HPE), Matt Drozt (HPE)

November 16th, 2025

Comparing Distributed-Memory Programming Frameworks


We see a need to compare distributed-memory programming frameworks

- To help programmers choose appropriate tools for their tasks
- To show which problems each framework excels at
- To compare both productivity and performance

Inspiration: the Computer Programming Languages Benchmark Game

- Anybody can contribute improved programs in any of the languages
- Website includes regularly updated performance results

To keep the scope manageable, we focused on a single benchmark



<https://benchmarksgame-team.debian.net/benchmarksgame/index.html>


New benchmarks & implementations are welcome in our repository!

- <https://github.com/hpc-ai-adv-dev/distributed-lsb/>




Experiment Overview

- Investigated five distributed programming frameworks on both **HPE Cray EX** (EX) and **InfiniBand** (IB) systems:
 - **Chapel** (with and without aggregation)
 - **OpenSHMEM** (with and without Conveyors for message aggregation)
 - **MPI**
 - **Lamellar** (currently only IB support)
- Implemented a communication intensive algorithm: **Distributed Least Significant Digit-First Radix Sort** (LSD Sort)



Methodology and Results

- **Development constraint:**
~1 day of implementation effort per framework by an experienced developer
- **Tuning:** Tested several configurations to identify optimal launch parameters
- **Benchmark:** Measured time to sort datasets of random **128-bit key/value pairs**



Source Lines of Code and Performance Details

Framework	Version	Lines of Code	EX Sorting Performance ¹	IB Sorting Performance ²
--- Compared LSD Sorting Implementations ---				
Chapel	Simple, Fine-Grained	★ 110	10,455	182
Chapel	Aggregated	113	★ 16,782	★ 2,519
MPI	(AllToAll)	● 410	7,823	1,018
OpenSHMEM	Simple, Fine-Grained	291	● 3,786	● 48
Conveyors	Aggregated OpenSHMEM	330	14,687	1,323
Lamellar	Unsafe Array	185	--	475
Lamellar	Atomic Array	175	--	468

1. Sorting 2^{37} elements using 128 Nodes (M elements sorted / second)

2. Sorting 2^{34} elements using 32 Nodes (M elements sorted / second)



The 4-Line Change in the Chapel Version

Chapel saw a **1.6-13.8x** improvement ...
with a small four-line change

```
// coforall task begin { ...
    var tD = calcBlock(task, 1D.low, 1D.high);
    // calc new position and put data there in temp
    {

        for i in tD {
            const ref tempi = temp.localAccess[i];
            const key = comparator.key(tempi);
            var bucket = getDigit(key, rshift, last, negs);
            var pos = taskBucketPos[bucket];
            taskBucketPos[bucket] += 1;
            a[pos] = tempi;
        }
    }
} //coforall task
```



The 4-Line Change in the Chapel Version

```
+ use CopyAggregation;

Chapel saw a 1.6-13.8x improvement ... // coforall task begin { ...
with a small four-line change
+           var tD = calcBlock(task, 1D.low, 1D.high);
+           // calc new position and put data there in temp
+           {
+               var aggregator = new DstAggregator(t);
+               for i in tD {
+                   const ref tempi = temp.localAccess[i];
+                   const key = comparator.key(tempi);
+                   var bucket = getDigit(key, rshift, last, negs);
+                   var pos = taskBucketPos[bucket];
+                   taskBucketPos[bucket] += 1;
+                   aggregator.copy(a[pos], tempi); // a[pos] = tempi;
+               }
+               aggregator.flush();
+           }
} //coforall task
```




What Changed From OpenSHMEM to Conveyors? One part:

```
for (int64_t i = 0; i < locN; i++) {  
  
    SortElement elt = localPart[i];  
  
    int bucket = getBucket(elt, digit);  
  
    int64_t &next = (*starts)[bucket];  
  
    int64_t dstGlobalIdx = next;  
  
    // store 'elt' into 'dstGlobalIdx'  
  
    auto dst = B.globalIdxToLocalIdx(dstGlobalIdx);  
  
    shmem_putmem(GB + dst.locIdx, &elt, sizeof(SortElement), dst.rank);  
  
    next += 1;  
}  
  
convey_begin(request, sizeof(IdxSortElement), alignof(IdxSortElement));  
  
int64_t i = 0;  
  
while (convey_advance(request, i == locN)) {  
  
    for (; i < locN; i++) {  
  
        SortElement elt = localPart[i];  
  
        int bucket = getBucket(elt, digit);  
  
        int64_t &next = (*starts)[bucket];  
  
        int64_t dstGlobalIdx = next;  
  
        // store 'elt' into 'dstGlobalIdx'  
  
        auto dst = B.globalIdxToLocalIdx(dstGlobalIdx);  
  
        IdxSortElement payload = { .locIdx = dst.locIdx, .value = elt };  
  
        if (!convey_push(request, &payload, dst.rank)) break;  
  
        next += 1;  
    }  
  
    IdxSortElement* local;  
  
    while((local = (IdxSortElement*)convey_apull(request, NULL)) != NULL) {  
  
        GB[local->locIdx] = local->value;  
    }  
}  
  
convey_reset(request);
```

Adding Conveyors
involves control flow
adjustment

Weak-Scaling Results Across Frameworks



- Problem Sizes:
 - EX: 2^{30} elements per node
 - IB: 2^{29} elements per node
- All the LSD sort implementations showed good weak scaling
- The aggregated Chapel implementation achieved the **highest performance** on both systems

How Do Other Distributed Sort Implementations Compare?

Framework	Version	Lines of Code	EX Sorting Performance ¹	IB Sorting Performance ²
--- Compared LSD Sorting Implementations ---				
Chapel	Simple, Fine-Grained	★ 110	10,455	182
Chapel	Aggregated	113	★ 16,782	★ 2,519
MPI	(AllToAll)	● 410	7,823	1,018
OpenSHMEM	Simple, Fine-Grained	291	● 3,786	● 48
Conveyors	Aggregated OpenSHMEM	330	14,687	1,323
Lamellar	Unsafe Array	185	--	475
Lamellar	Atomic Array	175	--	468

1. Sorting 2^{37} elements using 128 Nodes (M elements sorted / second)

2. Sorting 2^{34} elements using 32 Nodes (M elements sorted / second)




How Do Other Distributed Sort Implementations Compare?

Framework	Version	Lines of Code	EX Sorting Performance ¹	IB Sorting Performance ²
--- Compared LSD Sorting Implementations ---				
Chapel	Simple, Fine-Grained	⭐ 110	10,455	182
Chapel	Aggregated	113	⭐ 16,782	⭐ 2,519
MPI	(AllToAll)	ⓧ 410	7,823	1,018
OpenSHMEM	Simple, Fine-Grained	291	ⓧ 3,786	ⓧ 48
Conveyors	Aggregated OpenSHMEM	330	14,687	1,323
Lamellar	Unsafe Array	185	--	475
Lamellar	Atomic Array	175	--	468
--- Additional Sorting Implementations for Comparison ---				
Chapel	MSD Sort (Standard Library)	2,200	27,555	2,432
MPI	KaDiS AMS Sort	4,200	29,209	--

1. Sorting 2^{37} elements using 128 Nodes (M elements sorted / second)

2. Sorting 2^{34} elements using 32 Nodes (M elements sorted / second)


Comparing Performance on 128 HPE Cray EX Nodes



- Bars shown above the dashed line are comparable implementations of LSD sort
 - Chapel using aggregation:
 - ~**4x** faster than OpenSHMEM
 - **2x** faster than MPI
 - Chapel & Conveyors implementations benefit from:
 - message aggregation
 - one-sided communication
 - communication overlap
- Bars shown below the dashed line are other, more complex, distributed sorting algorithms for comparison



Comparing Performance and Productivity



- This scatterplot compares
 - time to sort 2^{37} elements using 128 nodes on EX
 - the number of source lines of code as counted by `cloc`
- Ideal: near origin in the lower left
- Of the LSD sort implementations, the aggregated Chapel version is the most performant while containing **~4x fewer lines of source code** as compared to MPI
- What made some implementations shorter?
 - native support for distributed arrays
 - built-in parallel prefix sum (aka scan)

A Qualitative Comparison of Examined Frameworks

MPI	Chapel	OpenSHMEM	Lamellar
<ul style="list-style-type: none">Pro: Industry standard, large developer communityCon: Missing distributed arrays, automatic aggregationCon: Often requires additional parallel strategies (OpenMP, CUDA, etc); we observed best performance with fewer ranks than cores	<ul style="list-style-type: none">Pro: Concisely create parallel tasks across cores and nodesPro: Aggregation and sorting in the standard library help productivity in distributed environmentsCon: Long compile times	<ul style="list-style-type: none">Pro: Very easy to installPro: Fast compile timesCon: Uses a custom collective allocator, limiting compatibility with C++ data types (e.g. std::vector)Con: Requires the use of other libraries (Conveyors) to provide aggregation with significant code changes	<ul style="list-style-type: none">Pro: Very easy to install as a Rust cratePro: Emphasis on memory safetyPro: Descriptive error messages help to catch bugs at compile timeCon: Currently only supports InfiniBand systemsCon: Long compile times (typical of Rust projects)Con: Not as mature

Contribute to this Study

- Repository can be found at:
<https://github.com/hpc-ai-adv-dev/distributed-lsb/>
- Contributions welcome and appreciated!
- We plan to make this a *living study repository* and update the performance results when new or improved LSD sort implementations are added



Thank You – Happy Sorting

Michael Ferguson – michael.ferguson@hpe.com

Ryan Friese – ryan.friese@pnnl.com

Shreyas Khandekar – shreyas.khandekar@hpe.com

Matt Drozt – drozt@hpe.com